

TENSORFLOW BEGINNER TUTORIAL

Sanket Shanbhag ¹

ABSTRACT

This document describes some basic concepts on how to use [tensorflow](#). Familiarity with machine learning concepts is assumed. We will be using Python3 along with tensorflow version 1.2.0.

INTRODUCTION

To do efficient numerical computing in Python, we typically use libraries like NumPy that do expensive operations such as matrix multiplication outside Python, using highly efficient code implemented in another language. Unfortunately, there can still be a lot of overhead from switching back to Python every operation. This overhead is especially bad if you want to run computations on GPUs or in a distributed manner, where there can be a high cost to transferring data.

TensorFlow also does its heavy lifting outside Python, but it takes things a step further to avoid this overhead. It does this by building a so called computational graph before-hand and performing all computations afterwards on this graph using highly efficient hardware specific instructions. Since GPU's are specifically designed for optimizing matrix multiplications, using TensorFlow allows us to harness the entire power of the GPU to speed up computations.

SETTING UP

Before we use Tensorflow, we must import it:

```
1 import tensorflow as tf
```

The best place to find more information on all the functions we will use is the [API docs](#).

The central unit of data in TensorFlow is the tensor. A tensor can simply be thought of as a multidimensional array. A tensor's rank is its number of dimensions. Here are some examples of tensors:

```
1 3 # a rank 0 tensor; this is a scalar with shape []
2 [1, 2, 3] # a rank 1 tensor; this is a vector with shape [3]
3 [[1, 2, 3], [4, 5, 6]] # a rank 2 tensor; a matrix with shape [2, 3]
4 [[[1, 2, 3]], [[7, 8, 9]]] # a rank 3 tensor with shape [2, 1, 3]
```

To create a computational graph, we create nodes and then run a session on these nodes to generate the output. Each node takes zero or more tensors as inputs and produces a tensor as an output.

INPUTS AND SESSIONS

Constants

One type of node is a constant. Like all TensorFlow constants, it takes no inputs, and it outputs a value it stores internally.

```
1 # Use dtype to optionally specify a type
2 node = tf.constant(42.0, dtype=tf.float32)
```

At this stage node is a tensor object that, when evaluated will hold the value 3.0. To actually evaluate this node, we have to run the computational graph in a session.

```
1 # Start a new Tensorflow session
2 sess = tf.Session()
3 print(sess.run([node])) # Prints 42
```

Placeholders

Constants are not that interesting, as they cannot be changed. To accept external inputs at run-time, we use placeholders.

```
1 a = tf.placeholder(tf.float32)
2 b = tf.placeholder(tf.float32)
3 adder_node = tf.add(a,b)
```

To give a value to the placeholder in the computational graph, use the `feed_dict` argument in the `run` function of the Session object to pass a python dictionary specifying the placeholders as key-value pairs.

```
1 print(sess.run(adder_node, feed_dict={a:18, b:24})) # Prints 42!
```

To assign placeholders of a higher rank, use the `shape` argument to specify the shape of the tensor. If any of the dimensions can be arbitrary, for example when using arbitrary number of training samples, you can use `None` instead.

```
1 # Takes a tensor of dimensions [None, 500]
2 a = tf.placeholder(tf.float32, shape=[None, 500])
3 # Tensor of dimensions [40, 50, 100]
4 b = tf.placeholder(tf.float32, shape=[40, 50, 100])
```

Variables

Variables allow us to add trainable parameters to a graph. A variable maintains state in the graph across calls to `run()`. You add a variable to the graph by constructing an instance of the class `Variable`. They are constructed with a type and initial value:

```
1 # Create some variables.
2 W = tf.Variable([.3], dtype=tf.float32)
3 # tf.random_normal(): Outputs random values from a normal
4 # distribution.
5 weights = tf.Variable(tf.random_normal([784, 200], stddev=0.35),
6                        name="weights")
```

```

7 # tf.zeros(): Creates a tensor with all elements set to zero
8 biases = tf.Variable(tf.zeros([200]), name="biases")

```

Like `tf.random_normal()` and `tf.zeros()` shown above, TensorFlow provides a collection of ops that produce tensors often used for [initialization from constants or random values](#).

Variables are not initialized when you call `tf.Variable`. To initialize all the variables in a TensorFlow program, you must explicitly call a special operation as follows:

```

1 sess.run(tf.global_variables_initializer())

```

SIMPLE LINEAR MODEL

In this section, we will create a simple model for classification on the easily available [MNIST](#) data-set and train it. Some familiarity with neural networks, activation functions and backpropagation are prerequisites for this section.

Downloading and Formatting Data

We will be using a single hidden layer of 500 nodes and an output layer of 10 classes.

```

1 import tensorflow as tf
2 from tensorflow.examples.tutorials.mnist import input_data
3 mnist = input_data.read_data_sets("./data/", one_hot = True)

```

This will download the MNIST data-set into a `data` folder in the current working directory. The data-set has been loaded as so-called One-Hot encoding. This means the labels have been converted from a single number to a vector whose length equals the number of possible classes. All elements of the vector are zero except for the `i`'th element which is one and means the class is `i`. For example, if the class value is 4, then it's one-hot encoded vector will be:

```

1 # 4th value is 1, everything else is 0
2 [0, 0, 0, 0, 1, 0, 0, 0, 0, 0]

```

Setting up Hyperparameters

```

1 # Hyperparameters
2 # 10 digits to identify
3 n_classes = 10
4 # Train 100 images at a time to avoid using up too much RAM
5 batch_size = 100
6 # We will be using a single hidden layer of 500 neurons
7 n_nodes_hl1 = 500
8 # The images are 28x28 pixels each
9 img_size_flat = 28 * 28

```

Setting up the model

To input images to our model, we will first flatten them into a single dimensional vector of length `img_size_flat`, and then feed this vector into our model. To do this, we will define a few placeholders:

```

1 # Input placeholder for flattened images
2 X = tf.placeholder('float', shape=[None, img_size_flat])
3 # Input vector for true class labels.
4 y = tf.placeholder('float')

```

Now we define the weights and biases for the hidden layer and output layer:

```

1 # Hidden layer
2 hidden_layer_weights = tf.Variable(tf.random_normal([img_size_flat,
3                                                       n_nodes_hl1]))
4 hidden_layer_biases = tf.Variable(tf.random_normal([n_nodes_hl1]))
5 # Output layer
6 output_layer_weights = tf.Variable(tf.random_normal([n_nodes_hl1,
7                                                       n_classes]))
8 output_layer_biases = tf.Variable(tf.random_normal([n_classes]))

```

Connecting the Layers

We now define the relationship between our layers by matrix multiplying the data with the weights and adding the biases. We use the [ReLU](#) activation function. Other activation functions are also available in tensorflow and can be found [here](#).

```

1 l1 = tf.add(tf.matmul(X, hidden_layer_weights),
2             hidden_layer_biases)
3 # Using ReLu activation
4 l1 = tf.nn.relu(l1)
5 # You can use the + operator instead of using the tf.add() function
6 output = tf.matmul(l1, output_layer_weights) + output_layer_biases

```

Training

We now define a function to train this network. We will use the [GradientDescent](#) Optimizer to reduce the [mean squared error \(MSE\)](#) loss function. Other optimizers in tensorflow can be found [here](#).

```

1 def train_neural_network(x):
2     prediction = x
3     # Softmax function
4     smx = tf.nn.softmax_cross_entropy_with_logits(logits=prediction,
5                                                    labels=y)
6     # Define the cost function
7     cost = tf.reduce_mean(smx)
8     # Use the GradientDescentOptimizer with a learning rate of 0.5
9     # to minimize the cost
10    optimizer = tf.train.GradientDescentOptimizer(0.5).\
11        minimize(cost)
12    # Number of epochs to train for
13    hm_epochs = 20
14
15    # Start a new tensorflow session
16    with tf.Session() as sess:
17        # Initialize the global variables
18        sess.run(tf.global_variables_initializer())
19        # Run a loop for the total number of epochs
20        for epoch in range(hm_epochs):

```

```

21     epoch_loss = 0
22     # Divide the data set into batches of size batch_size
23     batchquot = int(mnist.train.num_examples / batch_size)
24
25     for _ in range(batchquot):
26         # Get a batch of images and labels
27         xt, yt = mnist.train.next_batch(batch_size)
28
29         # Run the optimizer to minimize the
30         # cost on the batch
31         _, c = sess.run([optimizer, cost],
32                         feed_dict={X:xt, y:yt})
33         # Add the cost to our epoch loss
34         epoch_loss += c
35
36     print('Epoch', epoch+1, 'completed out of', hm_epochs,
37           'loss:', epoch_loss)
38
39     # Caculate and print accuracy
40     correct = tf.equal(tf.argmax(prediction, 1),
41                       tf.argmax(y, 1))
42     accuracy = tf.reduce_mean(tf.cast(correct, 'float'))
43     print('Accuracy:', accuracy.eval({X:mnist.test.images,
44                                       y:mnist.test.labels}))
44

```

To begin training, we simply pass our model to the function:

```

1 train_neural_network(output)

```

Result

Running the above network gives us the following results:

```

1 Epoch 1 completed out of 20 loss: 4112.72234179
2 Epoch 2 completed out of 20 loss: 292.313294291
3 Epoch 3 completed out of 20 loss: 179.839554987
4 Epoch 4 completed out of 20 loss: 126.560542699
5 Epoch 5 completed out of 20 loss: 97.231446553
6 Epoch 6 completed out of 20 loss: 76.5352744549
7 Epoch 7 completed out of 20 loss: 64.625726237
8 Epoch 8 completed out of 20 loss: 54.0737959952
9 Epoch 9 completed out of 20 loss: 47.0280316649
10 Epoch 10 completed out of 20 loss: 41.1886193645
11 Epoch 11 completed out of 20 loss: 36.654014512
12 Epoch 12 completed out of 20 loss: 32.950176964
13 Epoch 13 completed out of 20 loss: 29.3913420243
14 Epoch 14 completed out of 20 loss: 26.9070334777
15 Epoch 15 completed out of 20 loss: 24.5256814114
16 Epoch 16 completed out of 20 loss: 21.9067392419
17 Epoch 17 completed out of 20 loss: 20.2392465728
18 Epoch 18 completed out of 20 loss: 18.6161083714
19 Epoch 19 completed out of 20 loss: 17.0158358993
20 Epoch 20 completed out of 20 loss: 15.9119512606
21 Accuracy: 0.943

```

CONCLUSION

We learnt the basics of how to use tensorflow. It is important to get hands-on experience with TensorFlow in order to learn how to use it properly. Try changing the learning rate, batch size and the optimizers used and see how they impact speed and performance. Also try to increase the number of layers to see how it affects the accuracy.